


Remote Procedure Call

Outline

- Protocol Stack


- Presentation Formatting

RPC Timeline




RCP Components

- Protocol Stack
 - BLAST: fragments and reassembles large messages
 - CHAN: synchronizes request and reply messages
 - SELECT: dispatches request to the correct process
- Stubs



Bulk Transfer (BLAST)

- Unlike AAL and IP, tries to recover from lost fragments
- Strategy
 - selective retransmission
 - aka partial acknowledgements



BLAST Details

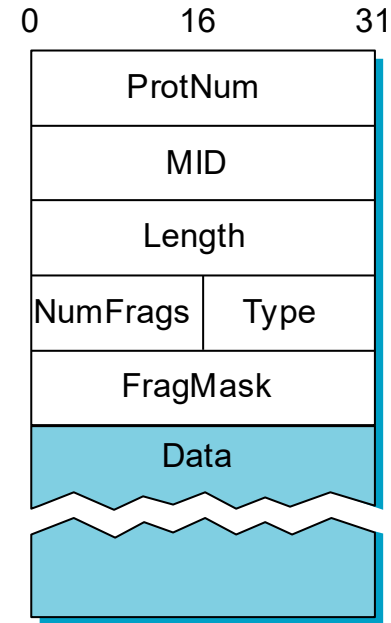
- Sender:
 - after sending all fragments, set timer DONE
 - if receive SRR, send missing fragments and reset DONE
 - if timer DONE expires, free fragments

BLAST Details (cont)

- Receiver:
 - when first fragments arrives, set timer LAST_FRAG
 - when all fragments present, reassemble and pass up
 - four exceptional conditions:
 - if last fragment arrives but message not complete
 - send SRR and set timer RETRY
 - if timer LAST_FRAG expires
 - send SRR and set timer RETRY
 - if timer RETRY expires for first or second time
 - send SRR and set timer RETRY
 - if timer RETRY expires a third time
 - give up and free partial message

BLAST Header Format


- MID must protect against wrap around
- TYPE = DATA or SRR
- NumFrag indicates number of fragments
- FragMask distinguishes among fragments
 - if Type=DATA, identifies this fragment
 - if Type=SRR, identifies missing fragments




Request/Reply (CHAN)

- Guarantees message delivery
- Synchronizes client with server
- Supports *at-most-once* semantics

Simple case



Implicit Acks



CHAN Details

- Lost message (request, reply, or ACK)
 - set RETRANSMIT timer
 - use message id (MID) field to distinguish
- Slow (long running) server
 - client periodically sends “are you alive” probe, or
 - server periodically sends “I’m alive” notice
- Want to support multiple outstanding calls
 - use channel id (CID) field to distinguish
- Machines crash and reboot
 - use boot id (BID) field to distinguish

CHAN Header Format

```
typedef struct {  
    u_short  Type;      /* REQ, REP, ACK, PROBE */  
    u_short  CID;       /* unique channel id */  
    int      MID;       /* unique message id */  
    int      BID;       /* unique boot id */  
    int      Length;    /* length of message */  
    int      ProtNum;   /* high-level protocol */  
} ChanHdr;
```

```
typedef struct {  
    u_char   type;       /* CLIENT or SERVER */  
    u_char   status;     /* BUSY or IDLE */  
    int      retries;    /* number of retries */  
    int      timeout;    /* timeout value */  
    XkReturn ret_val;    /* return value */  
    Msg      *request;   /* request message */  
    Msg      *reply;     /* reply message */  
    Semaphore reply_sem; /* client semaphore */  
    int      mid;        /* message id */  
    int      bid;        /* boot id */  
} ChanState;
```

Synchronous vs Asynchronous Protocols

- Asynchronous interface

`send(Protocol llp, Msg *message)`

`deliver(Protocol llp, Msg *message)`

- Synchronous interface


`call(Protocol llp, Msg *request, Msg *reply)`

`upcall(Protocol hlp, Msg *request, Msg *reply)`

- CHAN is a hybrid protocol
 - synchronous from above: **call**
 - asynchronous from below: **deliver**


Dispatcher (SELECT)

- Dispatch to appropriate procedure
- Synchronous counterpart to UDP
- Implement concurrency (open multiple CHANs)




- Address Space for Procedures
 - flat: unique id for each possible procedure
 - hierarchical: program + procedure number

Simple RPC Stack




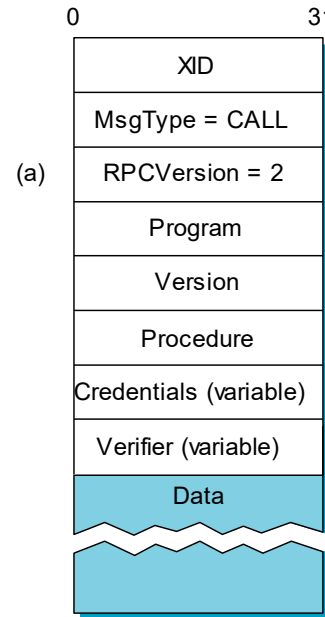
SunRPC

- IP implements BLAST-equivalent
 - except no selective retransmit
- SunRPC implements CHAN-equivalent
 - except not at-most-once
- UDP + SunRPC implement SELECT-equivalent
 - UDP dispatches to program (ports bound to programs)
 - SunRPC dispatches to procedure within program



SunRPC Header Format

- XID (transaction id) is similar to CHAN's MID
- Server does not remember last XID it serviced
- Problem if client retransmits request while reply is in transit



Presentation Formatting


- Marshalling (encoding) application data into messages
- Unmarshalling (decoding) messages into application data



- Data types we consider
 - integers
 - floats
 - strings
 - arrays
 - structs
- Types of data we do not consider
 - images
 - video
 - multimedia documents

Difficulties

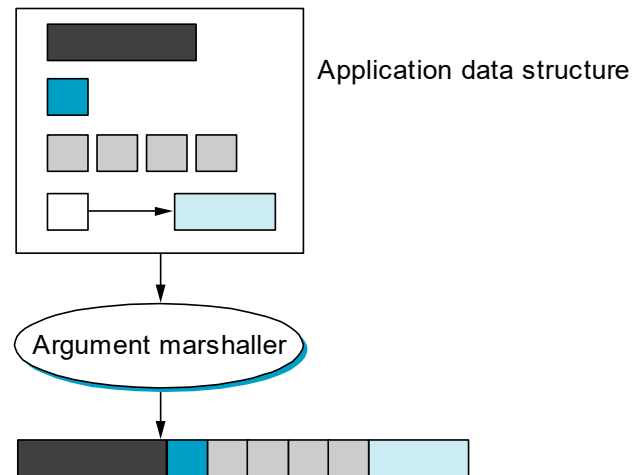
- Representation of base types
 - floating point: IEEE 754 versus non-standard
 - integer: big-endian versus little-endian (e.g., 34,677,374)



- Compiler layout of structures

Taxonomy

- Data types
 - base types (e.g., ints, floats); must convert
 - flat types (e.g., structures, arrays); must pack
 - complex types (e.g., pointers); must linearize



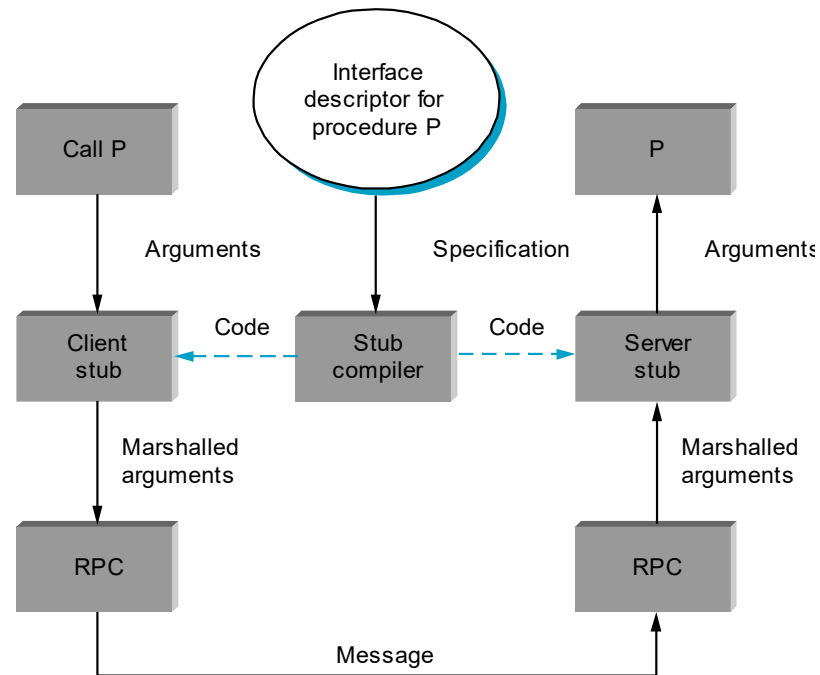
- Conversion Strategy
 - canonical intermediate form
 - receiver-makes-right (an $N \times N$ solution)

Taxonomy (cont)

- Tagged versus untagged data

type = INT	len = 4		value = 417892	
---------------	---------	--	----------------	--

- Stubs
 - compiled
 - interpreted




eXternal Data Representation (XDR)

- Defined by Sun for use with SunRPC
- C type system (without function pointers)
- Canonical intermediate form
- Untagged (except array length)
- Compiled stubs

```
#define MAXNAME 256;
#define MAXLIST 100;
```

```
struct item {
    int    count;
    char    name[MAXNAME] ;
    int    list[MAXLIST] ;
};
```

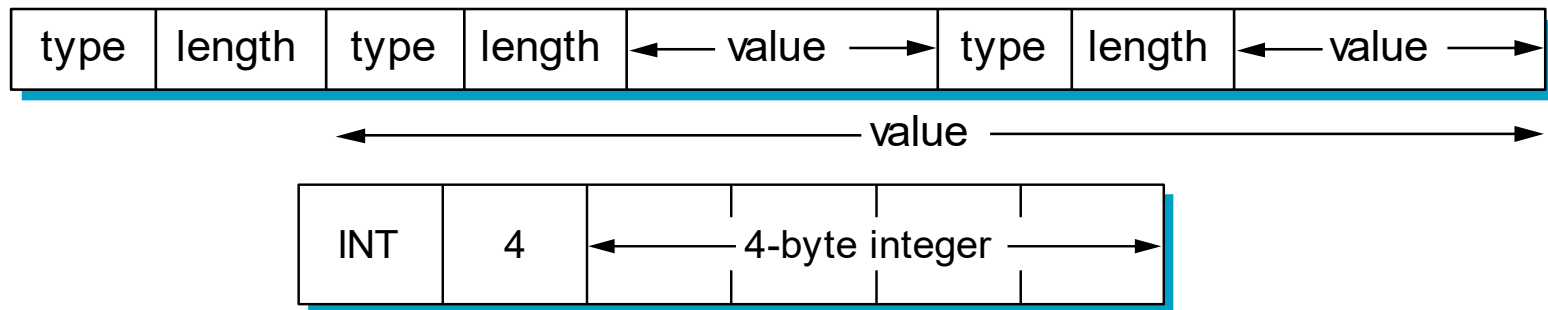
```
bool_t
xdr_item(XDR *xdrs, struct item *ptr)
{
    return(xdr_int(xdrs, &ptr->count) &&
        xdr_string(xdrs, &ptr->name, MAXNAME) &&
        xdr_array(xdrs, &ptr->list, &ptr->count,
            MAXLIST, sizeof(int), xdr_int));
}
```



Abstract Syntax Notation One (ASN-1)

- An ISO standard
- Essentially the C type system
- Canonical intermediate form
- Tagged
- Compiled or interpreted stubs
- BER: Basic Encoding Rules

(tag, length, value)



Network Data Representation (NDR)

- Defined by DCE
 - Essentially the C type system
 - Receiver-makes-right (architecture tag)
 - Individual data items untagged
 - Compiled stubs from IDL
 - 4-byte architecture tag
- IntegerRep
 - 0 = big-endian
 - 1 = little-endian
 - CharRep
 - 0 = ASCII
 - 1 = EBCDIC
 - FloatRep
 - 0 = IEEE 754
 - 1 = VAX
 - 2 = Cray
 - 3 = IBM

